Active Maintenance Report: schwarz_coppock_2022

Schwarz and Coppock (2022, Journal of Politics): What Have We Learned about Gender from Candidate Choice Experiments?

2026-08-01

Table of contents

1 Summary **2**
- 1.1 Does the deposited archive run? . . . 3
- 1.2 Does the maintained rewrite reproduce the paper? . . . 3

2 Paper overview **4**

3 Original archive reproducibility **5**

4 Errata **6**

5 Ground truth **7**

6 Every published number **9**

7 Maintained rewrite **11**

8 Figures **13**

9 Rewrite verification **18**

10 R environment **19**

Drafted by Claude Opus 5 under the supervision of Alex Coppock.

This repository holds the actively maintained replication code for Schwarz and Coppock (2022), together with the reproducibility report that documents what the original archive did and did not do. It is part of a program applying the maintenance proposal in Peer, Orr and Coppock

(2021, *PS: Political Science & Politics*, doi 10.1017/S1049096521000366) to a set of published archives.

Article	10.1086/716290
Replication archive	10.7910/DVN/R27ULT

The data are not redistributed here. The deposit is 0.5 MB across 12 files and lives at Harvard Dataverse, which is the only copy this repository points at. `download_original.R` fetches it and verifies every file; `original_manifest.csv` pins the file identifiers, sizes and checksums, so the exact bytes this code was written against are recorded in version control even though the bytes themselves are not.

Repository layout. `maintained/` is the maintained rewrite: one script per published table or figure, writing to `output/`, which is committed so a reader can compare a fresh run against it without downloading anything. `ground_truth/` ties every published number to the code that produces it. `original/` is created by the download script and is deliberately absent from the repository. This file is the reproducibility report, also available as a PDF in `report/`.

License. CC0 1.0 Universal, matching the terms of the deposit this repository maintains. See `LICENSE`.

To reproduce. Clone or download the repository, open `schwarz_coppock_2022.Rproj`, and run:

```
source("run_all.R")
```

That fetches the deposit, verifies its 12 files, and produces every table and figure into `maintained/output/`. Required packages: tidyverse, metafor, knitr, kableExtra, here. Paths resolve through `here`, so nothing depends on the working directory. The full run takes about six seconds. A successful run overwrites `maintained/output/`, which is committed: **`git diff` on that folder is the reproduction check.**

1 Summary

Two questions, answered before the detail.

1.1 Does the deposited archive run?

Yes, completely. All five analysis scripts execute without error on a current R installation. There are no `setwd()` calls, no hardcoded absolute paths and no functions called before they are defined, and every package the scripts name still installs from CRAN, `rmeta` and `ggstance` included. Nothing in the deposit needs a substitution to run.

What the archive does not do is write anything. All eight `ggsave()` calls, all three `print.xtable()` calls and the single `sink()` block are commented out in the deposited files, so a clean run leaves nothing behind but a stray `Rplots.pdf` from the plots being printed to a device. There is nothing to diff a rerun against, which is why the maintained rewrite writes its outputs to disk.

1.2 Does the maintained rewrite reproduce the paper?

Mostly, and the exceptions are mostly in the article rather than in the code. Every number the article and its appendix print has been extracted by hand into `ground_truth/published_claims.csv`, 2,624 of them, and 2,587 reproduce exactly at the precision the page uses. 15 disagree and 22 have no counterpart in the deposit at all.

Most of that total is float cells that a reader would not think of as a table. The four forest plots print an estimate and a standard error beside **every** row, not only beside the pooled summary, which is 472 published numbers; the 44 per-study appendix figures do the same, which is another 1,888; and Table 1 prints subject and rating counts for each of the 67 studies, which is 136 more. All of those reproduce but two.

The disagreements, none of which changes a substantive conclusion:

Table 1: Published numbers the deposit does not reproduce.

Claim	Article	Deposit	Where the fault lies
figure_3 :: female :: sen_2 :: estimate	6.0	5.9	archive
figure_3 :: male :: sen_2 :: estimate	0.5	0.4	archive
table_2 :: trimmed :: label_percent	90	91	unresolved
theory :: share_positive_fraction	0.75	0.701	paper_internal
design :: n_papers	49	50	unresolved
results :: n_positive_significant	23	29	unresolved
results :: n_negative	18	19	paper_internal
results :: n_negative_significant	11	6	unresolved
results :: trim_percent	5	3.0 and 6.0	unresolved
results :: n_resp_gender_studies	36	37	paper_internal
discussion :: share_positive_fraction	0.75	0.701	paper_internal
table_a1 :: n_studies_listed	43	44	paper_internal
appendix_b :: n_digitally_measured_experiments	9	10	unresolved
appendix_b :: n_digitally_measured_papers	7	8	unresolved
figure_a18 :: printed_estimates	0	10	paper_internal

6 are the article contradicting its own data. Five of those are in `errata.qmd`, which renders to `schwarz_coppock_2022_errata.pdf` and sets out 11 numbered items, five of them computed from the deposit at render time and the rest corrections of references, wording, and captions; the sixth is the row count of appendix Table A.1, which is a defect in a table rather than in a sentence and is described below. 2 are the deposit disagreeing with the page. The remaining 7 are unresolved, in the strict sense that the cause cannot be established from the deposit and guessing at one would be worse than saying so.

One published figure prints no data at all. Appendix Figure A.18 is Costa (2020), whose randomized attribute levels are whole tweets. At the width the appendix uses, those row labels push the plotting panel, the points, the intervals, the estimate labels and the axis off the right edge of the page, so the published figure shows a column of truncated sentences, an axis title cut off after four letters, and none of its ten estimates. The deposited code reproduces the defect exactly, which is why the maintained `figure_a18_costa_1.pdf` looks the same; `figure_a18_costa_1_legible.pdf` is the same ten estimates redrawn so they can be read, and it is what the errata reprints.

The significance counts reproduce under no rule. The article reports that of the 67 estimates, "47 are positive (23 significant) and 18 are negative (11 significant)." The rule the deposited script itself applies, a confidence interval excluding zero, gives 29 and 6. Testing p below 0.05 gives 19 and 2, and cannot be what was done in any case, because 19 of the 67 rows carry no p-value at all. Nor is this a matter of finding the right threshold: both counts rise as a threshold is relaxed, so no threshold produces 23 and 11 together. The threshold that puts 23 estimates in the positive column leaves 5 in the negative one, and the threshold that reaches 11 in the negative column leaves 43 in the positive one. No corrected value can be stated, so none is; this is recorded as unresolved rather than as an erratum.

Every estimate, standard error, interval and p-value the article reports in prose reproduces exactly. What fails is a set of hand-counted descriptions of those estimates, plus two cells of one appendix figure.

2 Paper overview

Citation: Schwarz, S. and Coppock, A. (2022). "What Have We Learned about Gender from Candidate Choice Experiments? A Meta-Analysis of Sixty-Seven Factorial Survey Experiments." *The Journal of Politics*, 84(2), 655-668. DOI: 10.1086/716290

Summary: Sixty-seven conjoint and factorial candidate choice survey experiments, fielded between 1982 and 2021 in more than twenty countries, all randomized candidate gender alongside other candidate attributes. Each study's treatment effect of being described as a woman is re-estimated on a common footing and the 67 estimates are pooled by random-effects meta-analysis. The pooled effect is a gain of 1.8 percentage points of vote share for women candidates, with a confidence interval from 1.1 to 2.5 points. The effect is larger for white

than for Black candidates in the US context (2.2 against 0.9 points), larger among women respondents than men (3.0 against 0.9), and reverses sign among Republicans (-1.4) while running positive among Democrats (4.2) and Independents (3.3).

What the deposit contains: five pre-computed summary datasets and five analysis scripts. The individual-level data belong to the authors of the 67 original studies and are not redistributed, so the deposit begins one level up, at estimated ATEs and CATEs with their standard errors. Every number in the article is a function of those five files.

3 Original archive reproducibility

Table 2: The deposited scripts on a current R installation.

Script	Status on current R	Resolution
helpers.R	Clean (sourced by all five scripts)	No changes required
meta_analyze_sates.R	Runs; writes nothing, every output call is commented out	No changes required
meta_analyze_cates_cand_race.R	Runs; writes nothing	No changes required
meta_analyze_cates_resp_gender.R	Runs; writes nothing	No changes required
meta_analyze_cates_resp_pid.R	Runs; writes nothing. Reports the Independent-Republican p-value from the Democrat-Republican statistic	Corrected in the rewrite; see Errata
meta_analyze_cates_cand_attributes.R	Runs; writes nothing. Its per-study figure loop is written but its writes are commented out	Restored in the rewrite as 44 written figures

Nothing here needs fixing to make the archive run. The deprecations are all warnings rather than errors: `geom_errorbarh()` and the `%+%` operator are deprecated in ggplot2 4.0.0 but still work, `reshape2::melt()` and `dcast()` are superseded, and `dplyr::do()` has been superseded since 1.0.0. What is wrong with this archive is the p-value bug documented in the Errata, which is a logic error rather than a run failure.

`rmeta` and `ggstance` are both still installable. The maintained rewrite replaces `rmeta::meta.summaries()` with `metafor::rma()` on maintenance grounds rather than availability grounds: `rmeta` has not been updated since March 2018. The two implementations of DerSimonian-Laird agree on every pooled estimate, standard error and interval bound in the paper to more than ten significant figures. `ggstance::position_dodgev()` resolves without incident from ggstance 0.3.7, which `coefplot` pulls in. It is dropped from the rewrite for a different reason: it is inert. Building the layer with and without it puts every point at the same coordinates, because every y value it is asked to dodge is already unique, 954 distinct

labels across 954 rows in Figure A.1 and no repeated label within any study in Figures A.2 through A.45.

What the archive does not produce. Every write in the deposit is commented out: eight `ggsave()` calls, three `print.xtable()` calls and one `sink()` block. Two of the commented `ggsave()` paths point at a directory that was never part of the deposit (`04_Drafts/current_draft/figures/`). The consequence is that the deposit cannot be checked against itself. There is no committed output to diff a rerun against, and a reader who runs all five scripts sees console output and a `Rplots.pdf` and nothing else. Restoring those writes, so that a rerun can be compared against a committed result, is the single largest change the maintained rewrite makes.

4 Errata

This section documents corrections the maintained rewrite makes to the deposited code. Errors in the article itself are a separate document: `errata.qmd`, which renders to `schwarz_coppock_2022_errata.pdf`, sets out the 11 corrections one by one, with each of the five numeric ones computed from the deposit at render time. Of the three bibliographic entries, one comes from a sweep of all 119 printed references against Crossref, which drew a single flag, and it is real: the entry for the Clayton and Nyhan (2020) manuscript naming two of its six authors. The manuscript at the URL that entry itself prints is by Carey, Clayton, Helmke, Nyhan, Sanders and Stokes, so the short form the article uses as a study label in Table 1, the appendix study table and Figures A.16 and A.17 should read Carey et al. (2020). No estimate or sample size is affected.

The Independent-versus-Republican p-value is computed from the wrong statistic. In `meta_analyze_cates_resp_pid.R`, the block that tests the two differences between pooled partisan CATEs computes the second p-value from the first test's statistic:

```
# Archive:
I_R_statistic = I_R_estimate / I_R_std.error,
I_R_p.value = 2 * pnorm(abs(D_R_statistic), lower.tail = FALSE)

# Rewrite:
I_R_p.value = 2 * pnorm(abs(I_R_statistic), lower.tail = FALSE)
```

`I_R_statistic` is computed on the line above and then never used, so the reported Independent-Republican p-value is numerically identical to the Democrat-Republican one.

Quantity	Published	Original archive	Maintained rewrite
Independent minus Republican, p	not reported	2.21e-10 (equals the D-R p-value)	1.06e-07

The article states only that "the Republican estimate is statistically significantly different from the Democratic and Independent estimates," which is true of both the wrong value and the right one. No published number changes.

Not corrected: the trimming rule. Table 2's last-but-one row is labelled "Trimmed (middle 90% of estimates)" and the text says the row drops "the bottom and top 5% of point estimates." The code keeps estimates above the 2nd percentile and below the 95th, which drops the bottom 3.0 per cent and the top 6.0 per cent and retains 61 of 67 studies, or 91 per cent rather than 90. Correcting 0.02 to 0.05 would be a plausible reading of a typo, but it is a change to a sample definition rather than to a mechanical error, so the rewrite reproduces the published row as computed and reports the alternative separately in `output/table_2_trim_sensitivity.csv`. The symmetric rule the text describes gives 0.020 over 59 studies against the published 0.018 over 61, so the robustness claim the row is making survives either reading.

Not corrected: the year threshold. The subset labelled "Post-2014 studies" is selected by `year > 2010`. The labels are nonetheless accurate, because the sample contains no study fielded between 1998 and 2014, a gap the article notes in its fourth footnote, so every threshold in that window produces the same two groups of 63 and 4.

5 Ground truth

Table 3: Ground truth: 100 recorded quantities, by where the article states them.

Table or figure	Reproduces	Not stated in the article	Does not reproduce
figure_1	5	0	0
figure_2	6	5	0
figure_3	6	5	1
figure_4	7	8	0
figure_a1	3	0	0
table_1	3	1	0
table_2	40	1	0
table_a1	0	0	1
text	2	2	4

`ground_truth/schwarz_coppock_2022_ground_truth.csv` records 100 quantities and answers a question about the deposited code: `value_script` is what the deposited scripts

produce through `rmeta`, `value_rewrite` is what `maintained/` produces through `metafor`, and `value_paper` is filled only by reading the article. Where the article does not state a quantity the cell is empty and `match` is `NA`, which is the case for every confidence interval bound in the four figures: the forest plots draw the intervals but label only the estimate and its standard error.

The article's own side of the comparison is exhaustive and lives separately, in the next section.

Table 4: The 6 recorded quantities that do not reproduce.

Table or figure	Claim	Article	Deposit	Note
text	n_negative	18	19	Text: 18 are negative. The depos has 19 negative, one exactly zero and 47 positive, which sums to 6 the published 47 plus 18 sums to
text	n_pos_and_sig_ci_rule	23	29	Text: 47 estimates are positive (significant). The deposited script and the rewrite both give 29 und the interval rule the deposit appl and no rule the deposit supports gives 23. Both counts rise as a threshold is relaxed, so no thresh gives 23 and 11 at once: reaching negative would leave 44 positive. The count is stated without a hed so it is compared as printed.
text	n_neg_and_sig_ci_rule	11	6	Text: 18 are negative (11 significant). The deposited script and the rewrite both give 6 and rule reaches 11; see the note on t positive count.
text	prop_positive	0.75	0.7014925373134329	Discussion states more than three-fourths of the studies show positive effect. value_paper reco that lower bound rather than an equality; 47 of 67 is 0.701, so the claim does not hold
figure_3	n_studies_in_text	36	37	Text: for the 36 experiments whe we were able to identify responde gender. The deposit and the paper's own figure both have 37

table_a1	n_studies_listed	43	44	Published Table A.1 lists 43 stud Simas and Murdoch (2019) has a CATE figure at A.44 but no Tab A.1 row

6 Every published number

The ground truth above was built around the article's floats. That misses the numbers a reader is most likely to quote, which sit in sentences, and it misses the numbers a figure prints on its own face. So the article and its appendix were read end to end and every numeric token extracted by hand into `ground_truth/published_claims.csv`: 2,624 of them, each classified by what would verify it.

Table 5: The 2,624 numbers the article and appendix print.

What would verify it	Claims
Pipeline: the deposit should produce it	2585
Descriptive: a claim about shape, sign or a count	13
Transcribed: copied from another paper	13
Definitional: a scale point or a design parameter	7
Structural: a float's own shape	6

Two instruments read those claims against the pipeline, by separate paths from the same outputs. `ground_truth/build_ground_truth.R` writes `ground_truth/published_claims_reproduction.csv`, which puts the pipeline's value beside every published one; `maintained/in_text_claims.R` prints each number beneath the sentence that states it, quoted verbatim. Where one parses the label a plot or a table actually draws, the other formats the estimate behind it. The build runs the claims file, counts what it prints against the extraction, and stops if the two instruments disagree anywhere, so neither can quietly stop checking.

Table 6: Published numbers against the deposit, by where the article prints them.

Where the article prints it	Claims	Reproduces	Does not	No counterpart
abstract	2	2	0	0
appendix_a	1	0	0	1
appendix_b	5	0	2	3
design	12	9	1	2
discussion	7	4	1	2
figure_1	137	137	0	0

figure_2	68	68	0	0
figure_3	152	150	2	0
figure_4	116	116	0	0
figure_a1	1	1	0	0
figure_a2 to figure_a45	1889	1888	1	0
intro	11	4	0	7
results	34	29	5	0
table_1	136	136	0	0
table_2	41	40	1	0
table_a1	1	0	1	0
theory	10	2	1	7
title	1	1	0	0

Every forest plot in the paper is a published table in disguise. Each prints an estimate and a standard error, in percentage points to one decimal, beside every study row and not only beside the pooled summary. In the four main figures that is 472 numbers, and in the 44 per-study appendix figures another 1,888. The four main figures are rasterized in the published PDF, so no text search reaches them; they were read off page renderings. All but two reproduce, and the two are both Sen (2017), USA, SSI 2 in Figure 3, where the page prints 6.0 and 0.5 and the deposited respondent-gender file gives 5.9 and 0.4. The deposited study-level file does carry a revised Sen SSI 2 estimate, which Figure 1 reproduces exactly, so the conditional-effects file is one revision behind the figure the article printed. Every one of the 1,888 appendix labels reproduces exactly.

The deposit has no counterpart for 22 published numbers, each named with its reason in `build_ground_truth.R`. Twelve are quoted from other people's papers and cannot drift when this pipeline moves. Two are theoretical reference points, a majority threshold and the midpoint of a vote margin. Three are the appendix B validation of the digital measurement procedure, whose average absolute error and correlation compare digital measurements against estimates computed from Kirkland and Coppock's raw data; the deposit carries neither, so those numbers are not reconstructible from it. What the deposit does record is which studies were measured that way, in a free-text note field: the note marks ten experiments in eight papers where the appendix says nine in seven.

Two further disagreements the float-shaped ground truth could not see. The design section says the 67 experiments come from 49 papers, where the deposit assigns 50 distinct bibliography keys to them. And Table 2's trimmed row, labelled "middle 90% of estimates" and described in the text as dropping "the bottom and top 5% of point estimates," in fact drops the bottom 3.0 per cent and the top 6.0 per cent, retaining 61 of 67. Both are recorded as unresolved: the deposit is self-consistent in each case and cannot say which of the two readings was intended.

7 Maintained rewrite

Table 7: The maintained rewrite.

Script	Output
helpers.R	Packages, re_meta(), diff_test(), label formatting
table_1_study_manifest.R	table_1_study_manifest.csv
table_2_meta_robustness.R	table_2_meta_robustness.csv, table_2_trim_sensitivity.csv
table_a1_study_attributes.R	table_a1_study_attributes.csv
figure_1_sate_forest_plot.R	figure_1_sate_forest_plot.pdf and .png, figure_1_meta_summary.csv, figure_1_plotted_estimates.csv
figure_2_cate_candidate_race.R	figure_2_cate_candidate_race.pdf and .png, figure_2_meta_summary.csv, figure_2_diff_in_cates_cand_race.csv, figure_2_plotted_estimates.csv
figure_3_cate_respondent_gender.R	figure_3_cate_respondent_gender.pdf and .png, figure_3_meta_summary.csv, figure_3_diff_in_cates_resp_gender.csv, figure_3_plotted_estimates.csv
figure_4_cate_respondent_pid.R	figure_4_cate_respondent_pid.pdf and .png, figure_4_meta_summary.csv, figure_4_diff_in_cates_resp_pid.csv, figure_4_plotted_estimates.csv
figure_a1_cate_candidate_attributes.R	figure_a1_cate_candidate_attributes.pdf and .png, figure_a1_cate_summary.csv
figure_a2_a45_cate_by_study.R	44 per-study appendix figures, figure_a2_a45_index.csv
text_in_text_numbers.R	text_in_text_numbers.csv, text_study_level.csv
in_text_claims.R	Printed only: every published number beneath the sentence that states it

Substitutions.

Deposited	Maintained	Why
`rmeta::meta.summaries(method = "random")`	`metafor::rma(method = "DL")`	Same DerSimonian-Laird estimator; `rmeta` last updated 2018
`do(tidy(with(., ...)))`	`group_by() |> reframe(re_meta(...))`	`do()` superseded in dplyr 1.0.0
`reshape2::melt()` and `dcast()`	`pivot_wider()`, or dropped	The melt-dcast-clean_names round trip existed only to line two pooled rows up side by side
`geom_errorbarh(height = 0)`	`geom_linerange()`	Deprecated in ggplot2 4.0.0
`ggstance::position_dodgev()`	dropped	Still available, but inert: every y value it dodges is already unique
`ggplot() %+% data`	`ggplot(data)`	Deprecated in ggplot2 4.0.0
commented-out `ggsave()`, `print.xtable()`, `sink()`	`ggsave()` and `write_csv()` to `output/`	So that a rerun can be diffed against a committed result
`%>%`	`|>`	Native pipe

The estimator is named at every call site. `re_meta()` in `helpers.R` takes `method` as a required argument rather than defaulting to `DL`, so no script can pool without saying which estimator it pooled with. `table_2_meta_robustness.R` reports both: the DL columns reproduce the published table, and REML columns sit beside them as an addition rather than a correction. REML moves no estimate by more than 0.001 and widens most intervals slightly, which is what a different heterogeneity estimator on 67 studies should do.

Coverage the archive states but does not deliver. The deposit's README says `meta_analyze_cates_cand_attributes.R` produces "Figure A.1, Figures A.2 through A.45, Table A.1." It contains the code for all three, with the writes commented out. The rewrite restores the 44 per-study figures, which take four seconds of the run's six. Table A.1 is a partial exception: the published table pairs each study with a prose description of its randomized attributes ("Race, ballot length"), and the deposit carries only the variable names those descriptions were written from, so `table_a1_study_attributes.csv` lists the variable names and the counts instead. The study coverage is the checkable part, and it is where the 43-against-44 discrepancy shows up.

8 Figures

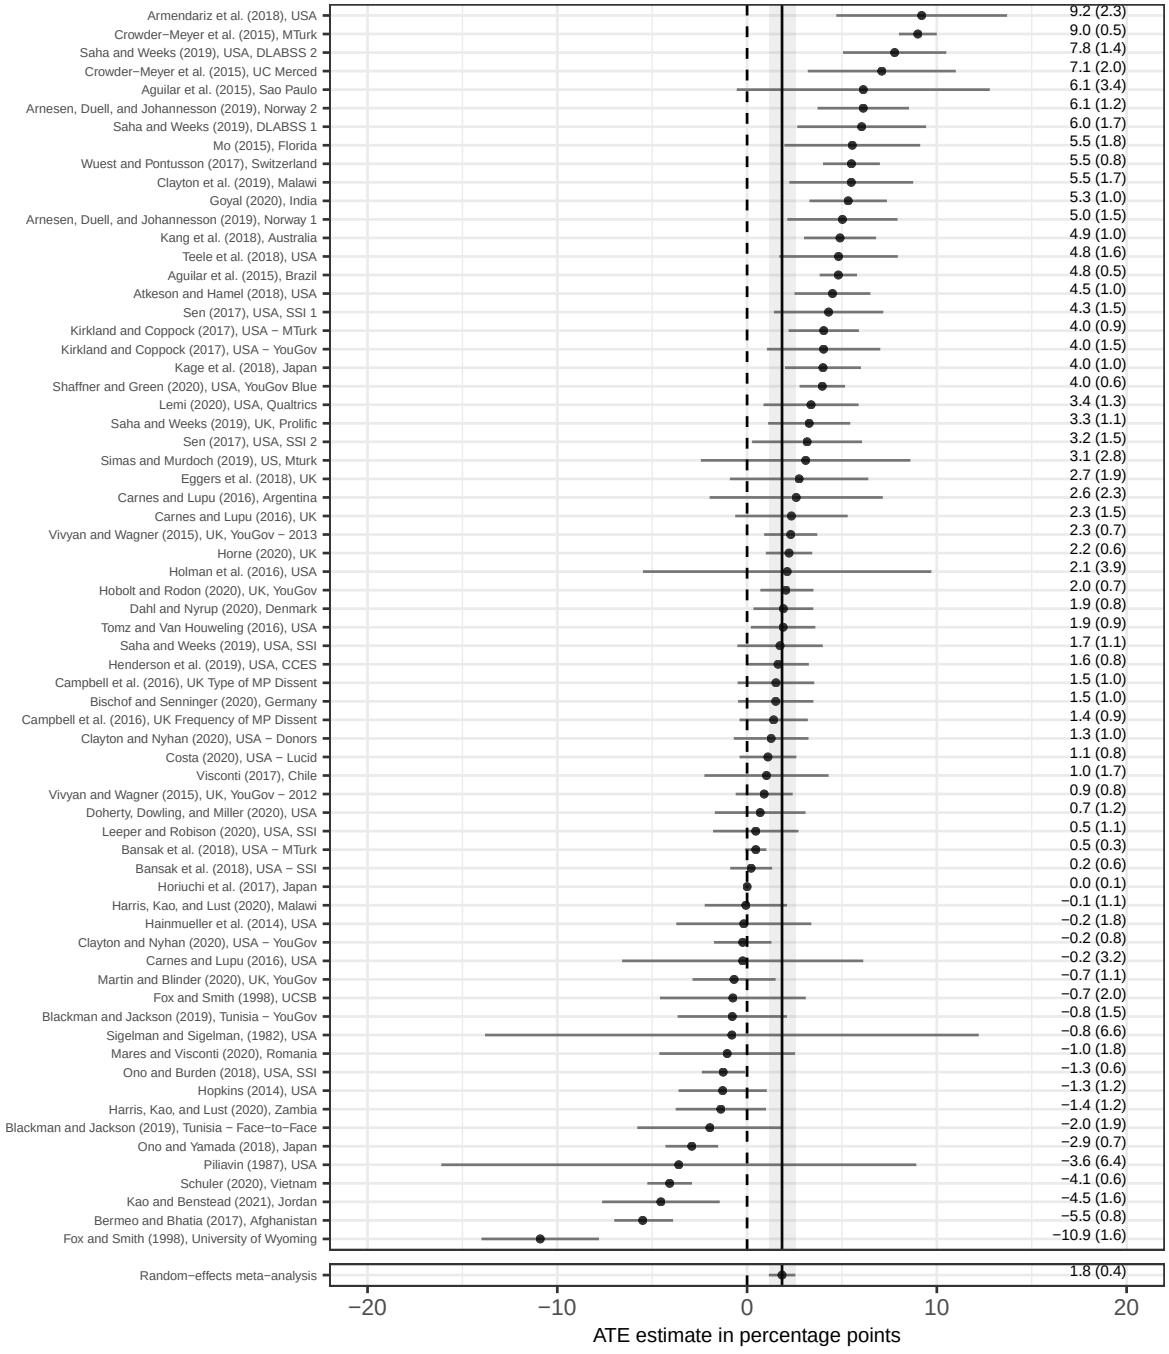


Figure 1: Figure 1. The 67 study-level ATEs and the pooled random-effects estimate.

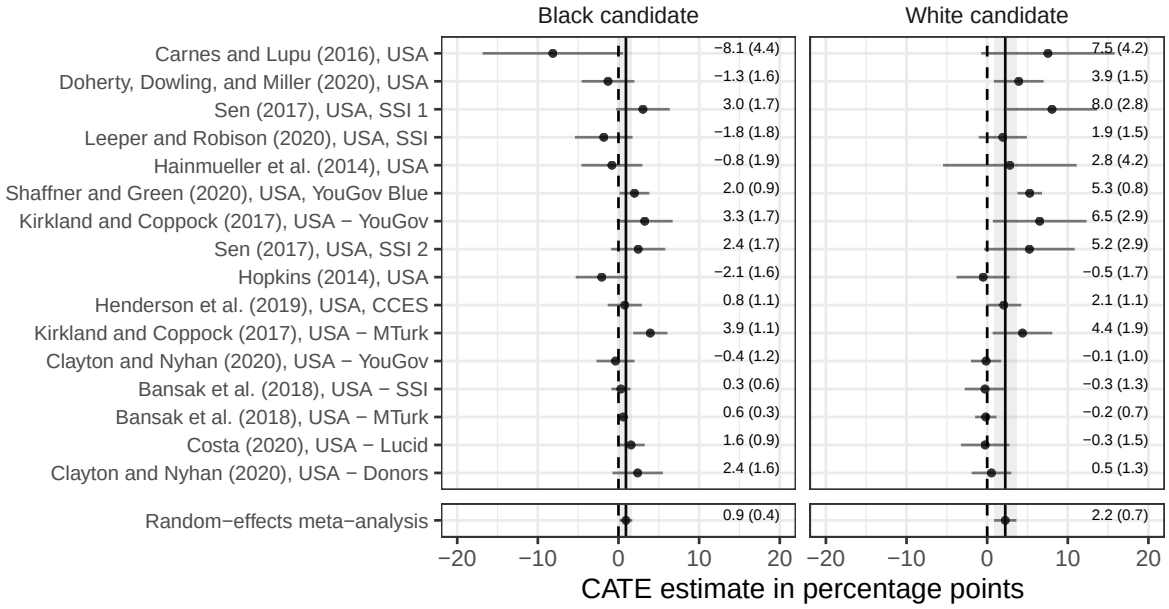


Figure 2: Figure 2. CATEs of candidate gender by candidate race.

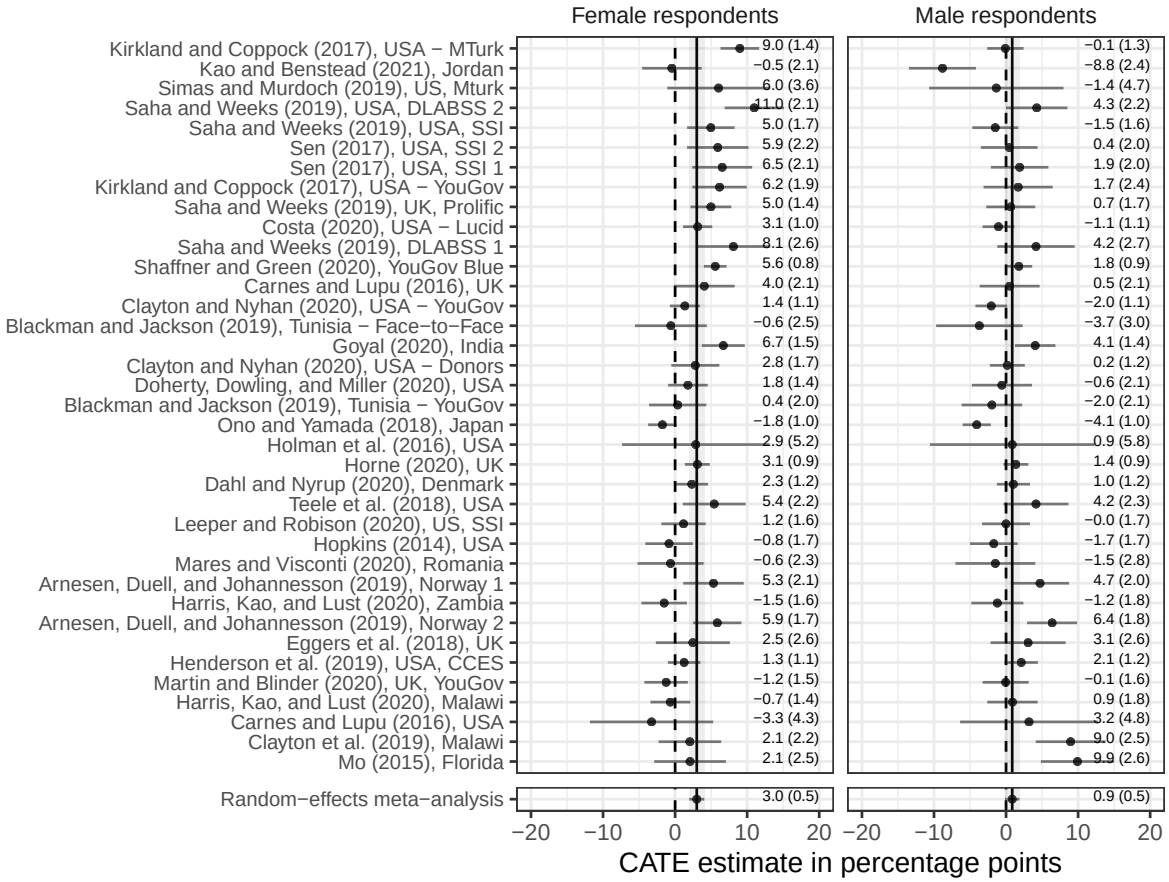


Figure 3: Figure 3. CATEs of candidate gender by respondent gender.

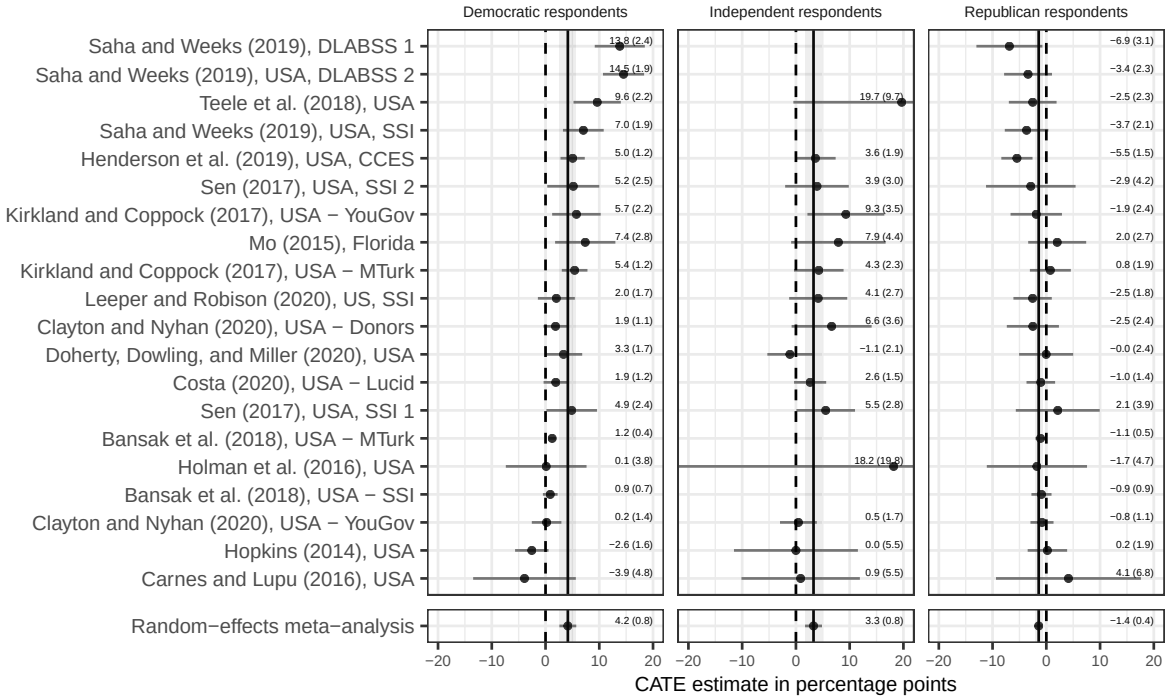


Figure 4: Figure 4. CATEs of candidate gender by respondent party identification.

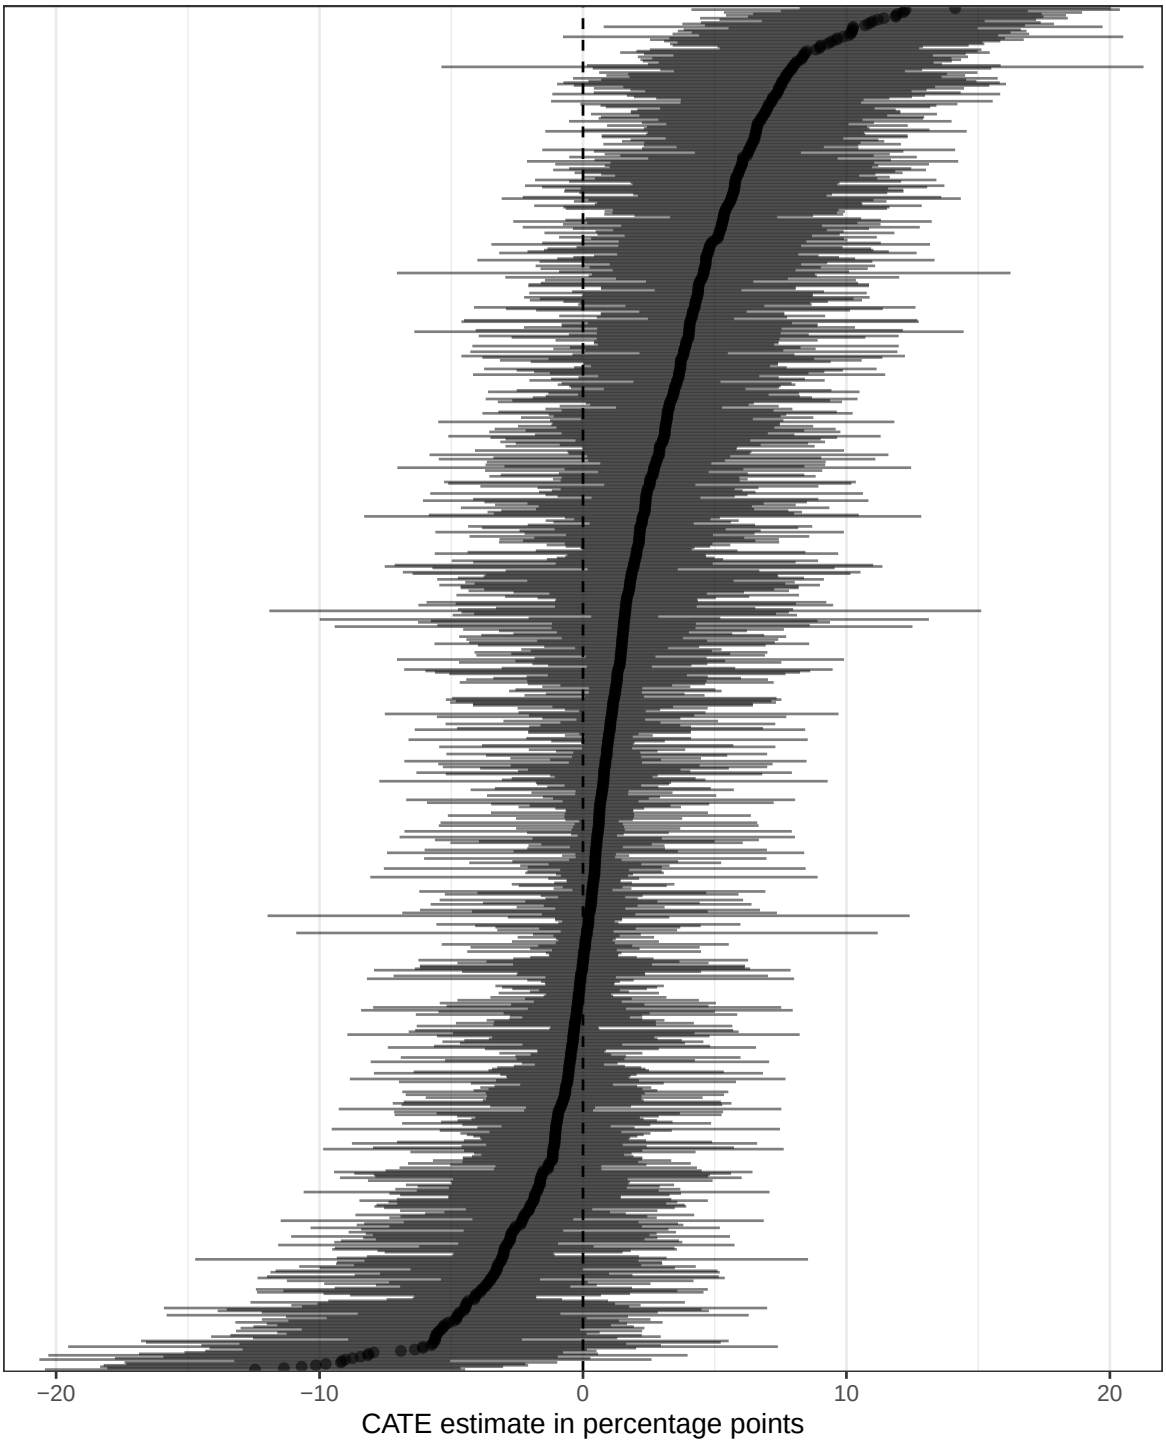


Figure 5: Figure A.1. All 954 CATEs of candidate gender conditional on the other randomized candidate attributes.

9 Rewrite verification

Table 8: Every quantity the article states, against the deposit and the rewrite (78 rows).

Table or figure	Claim	Article	Deposit	Rewrite	Deposit mat
table_1	n_studies	67	67	67	
table_1	n_subjects_total	120601	120601	120601	
table_1	n_ratings_total	774971	774971	774971	
figure_1	re_meta_estimate_pp	1.8	1.8421458705235139	1.842145870523515	
figure_1	re_meta_se_pp	0.4	0.354595479723342	0.354595479723342	
figure_1	re_meta_ci_low_pp	1.1	1.14715150118506	1.147151501185061	
figure_1	re_meta_ci_high_pp	2.5	2.537140239861967	2.537140239861968	
figure_1	n_studies	67	67	67	
text	n_positive	47	47	47	
text	n_negative	18	19	19	
text	n_exactly_zero	1	1	1	
text	n_pos_and_sig_ci_rule	23	29	29	
text	n_neg_and_sig_ci_rule	11	6	6	
text	prop_positive	0.75	0.7014925373134329	0.7014925373134329	
table_2	post_2014_studies_estimate	0.021	0.02076852592631652	0.02076852592631653	
table_2	post_2014_studies_std_error	0.004	0.00353901936870558	0.00353901936870559	
table_2	post_2014_studies_conf_low	0.014	0.01383217542306389	0.0138321754230639	
table_2	post_2014_studies_conf_high	0.028	0.02770487642956915	0.02770487642956917	
table_2	post_2014_studies_n_studies	63	63	63	
table_2	pre_1998_studies_estimate	-0.047	-0.04670204241768256	-0.04670204241768255	
table_2	pre_1998_studies_std_error	0.037	0.03660122845561661	0.0366012284556166	
table_2	pre_1998_studies_conf_low	-0.118	-0.11843913198061368	-0.11843913198061368	
table_2	pre_1998_studies_conf_high	0.025	0.02503504714524856	0.02503504714524856	
table_2	pre_1998_studies_n_studies	4	4	4	
table_2	convenience_student_sample_estimate	0.023	0.02297025791106249	0.02297025791106249	
table_2	convenience_student_sample_std_error	0.006	0.00566850898752522	0.00566850898752521	
table_2	convenience_student_sample_conf_low	0.012	0.01186018444947145	0.01186018444947146	
table_2	convenience_student_sample_conf_high	0.034	0.03408033137265354	0.03408033137265353	
table_2	convenience_student_sample_n_studies	31	31	31	
table_2	representative_sample_estimate	0.015	0.01493099437674829	0.01493099437674828	
table_2	representative_sample_std_error	0.005	0.00538611304870334	0.00538611304870334	
table_2	representative_sample_conf_low	0.004	0.00437440678462851	0.00437440678462851	
table_2	representative_sample_conf_high	0.025	0.02548758196886806	0.02548758196886806	
table_2	representative_sample_n_studies	36	36	36	
table_2	american_sample_estimate	0.023	0.02316935766935101	0.02316935766935102	
table_2	american_sample_std_error	0.006	0.00606511531362025	0.00606511531362025	
table_2	american_sample_conf_low	0.011	0.01128195009257296	0.01128195009257297	
table_2	american_sample_conf_high	0.035	0.03505676524612906	0.03505676524612908	
table_2	american_sample_n_studies	35	35	35	
table_2	non_american_sample_estimate	0.014	0.01374151072947667	0.01374151072947667	
table_2	non_american_sample_std_error	0.005	0.00484683076012992	0.00484683076012992	
table_2	non_american_sample_conf_low	0.004	0.00424189700046112	0.00424189700046112	
table_2	non_american_sample_conf_high	0.023	0.02324112445849222	0.02324112445849222	
table_2	non_american_sample_n_studies	32	32	32	
table_2	trimmed_middle_90_of_estimates_estimate	0.018	0.01760270448260373	0.01760270448260372	
table_2	trimmed_middle_90_of_estimates_std_error	0.003	0.00296095651140167	0.00296095651140167	

Table 8: Every quantity the article states, against the deposit and the rewrite (78 rows *(continued)*

Table or figure	Claim	Article	Deposit	Rewrite	Deposit mat
table_2	trimmed_middle_90_of_estimates_conf_low	0.012	0.01179933636046708	0.01179933636046707	
table_2	trimmed_middle_90_of_estimates_conf_high	0.023	0.0234060726047 4038	0.0234060726047 4037	
table_2	trimmed_middle_90_of_estimates_n_studies	61	61	61	
table_2	all_studies_estimate	0.018	0.01842145870523514	0.01842145870523515	
table_2	all_studies_std_error	0.004	0.00354595479723342	0.00354595479723342	
table_2	all_studies_conf_low	0.011	0.0114715150118506	0.01147151501185061	
table_2	all_studies_conf_high	0.025	0.02537140239861967	0.02537140239861968	
table_2	all_studies_n_studies	67	67	67	
figure_2	black_cand_estimate_pp	0.9	0.924095993624055	0.924095993624055	
figure_2	black_cand_se_pp	0.4	0.403949584944996	0.403949584944996	
figure_2	white_cand_estimate_pp	2.2	2.243283459448613	2.243283459448615	
figure_2	white_cand_se_pp	0.7	0.704882012460991	0.704882012460991	
figure_2	n_studies	16	16	16	
figure_2	diff_in_cates_pvalue	0.104	0.10442580239427804	0.1044258023942776	
figure_3	female_resp_estimate_pp	3.0	3.005785735908717	3.0057857359087192	
figure_3	female_resp_se_pp	0.5	0.532093754832003	0.532093754832003	
figure_3	male_resp_estimate_pp	0.9	0.85476192509097	0.85476192509097	
figure_3	male_resp_se_pp	0.5	0.509847917506916	0.509847917506916	
figure_3	n_studies_in_figure	37	37	37	
figure_3	n_studies_in_text	36	37	37	
figure_3	diff_in_cates_pvalue	0.004	0.0035127926564261	0.00351279265642609	
figure_4	dem_resp_estimate_pp	4.2	4.154259026957203	4.154259026957201	
figure_4	dem_resp_se_pp	0.8	0.799277010412441	0.799277010412441	
figure_4	ind_resp_estimate_pp	3.3	3.276867581970807	3.2768675819708095	
figure_4	ind_resp_se_pp	0.8	0.804774319348512	0.8047743193485131	
figure_4	rep_resp_estimate_pp	-1.4	-1.417655835707809	-1.417655835707809	
figure_4	rep_resp_se_pp	0.4	0.363468843734238	0.363468843734238	
figure_4	n_studies	20	20	20	
figure_a1	n_cates	954	954	954	
figure_a1	prop_cates_positive	0.5	0.69916142557652	0.69916142557652	
figure_a1	n_attribute_studies	44	44	44	
table_a1	n_studies_listed	43	44	44	

The rewrite reproduces the deposit exactly. Across all 100 recorded quantities the two agree to better than 1e-10 in every case but one, the Independent-Republican p-value corrected in the Errata above. Where the deposit fails to match the article, so does the rewrite, which is the intended behaviour: the failures are in counts the article states, not in anything the code computes.

10 R environment

Table 9: Packages loaded when this report was rendered.

Package	Version
dplyr	1.2.1
forcats	1.0.1
ggplot2	4.0.3
here	1.0.2
kableExtra	1.4.0
knitr	1.51
lubridate	1.9.5
purrr	1.2.2
readr	2.2.0
stringr	1.6.0
tibble	3.3.1
tidyr	1.3.2
tidyverse	2.0.0

Rendered under R version 4.6.1 (2026-06-24) on macOS Tahoe 26.6.2.